

Deep learning in computer vision and natural language processing

Yifeng Tao

School of Computer Science
Carnegie Mellon University

Slides adapted from Matt Gormley, Russ Salakhutdinov

Review


- Perceptron algorithm
- Multilayer perceptron and activation functions
- Backpropagation
- Momentum-based mini-batch gradient descent methods

Outline


- Regularization in neural networks – methods to prevent overfitting
- Widely used deep learning architecture in practice
 - CNN
 - RNN

Overfitting


- The model tries to learn too well the noise in training samples



Underfitting




Just right!



overfitting

Training Vs. Test Set Error




[Slide from <https://www.analyticsvidhya.com/blog/2018/04/fundamentals-deep-learning-regularization-techniques/>]

Model Selection

- Training Protocol:
 - Train your model on the **Training Set** $\mathcal{D}^{\text{train}}$
 - For model selection, use **Validation Set** $\mathcal{D}^{\text{valid}}$
 - Hyper-parameter search: hidden layer size, learning rate, number of iterations/epochs, etc.
 - Estimate generalization performance using the **Test Set** $\mathcal{D}^{\text{test}}$
- Generalization is the behavior of the model on **unseen examples**.

Regularization in Machine Learning

- Regularization penalizes the coefficients.
- In deep learning, it penalizes the weight matrices of the nodes.



Appropriate-fitting

[Slide from <https://www.analyticsvidhya.com/blog/2018/04/fundamentals-deep-learning-regularization-techniques/>]

Regularization in Deep Learning

- L2 & L1 regularization
- Dropout
- Data augmentation
- Early stopping
- Batch normalization

$$\arg \min_{\boldsymbol{\theta}} \frac{1}{T} \sum_t l(f(\mathbf{x}^{(t)}; \boldsymbol{\theta}), y^{(t)}) + \lambda \Omega(\boldsymbol{\theta})$$

L2 regularization:


$$\Omega(\boldsymbol{\theta}) = \sum_k \sum_i \sum_j \left(W_{i,j}^{(k)} \right)^2 = \sum_k \|\mathbf{W}^{(k)}\|_F^2$$

L1 regularization:

$$\Omega(\boldsymbol{\theta}) = \sum_k \sum_i \sum_j |W_{i,j}^{(k)}|$$

Dropout

- Produces very good results and is the most frequently used regularization technique in deep learning.
- Can be thought of as an ensemble technique.
 - Use random binary masks $m^{(k)}$
 - layer pre-activation for $k > 0$
$$\mathbf{a}^{(k)}(\mathbf{x}) = \mathbf{b}^{(k)} + \mathbf{W}^{(k)}\mathbf{h}^{(k-1)}(\mathbf{x})$$
 - hidden layer activation ($k=1$ to L):
$$\mathbf{h}^{(k)}(\mathbf{x}) = g(\mathbf{a}^{(k)}(\mathbf{x})) \odot m^{(k)}$$
 - Output activation ($k=L+1$)
$$\mathbf{h}^{(L+1)}(\mathbf{x}) = o(\mathbf{a}^{(L+1)}(\mathbf{x})) = \mathbf{f}(\mathbf{x})$$



[Slide from Russ Salakhutdinov et al.]

Dropout at Test Time

- At test time, we replace the masks by their expectation
 - This is simply the constant vector 0.5 if dropout probability is 0.5
 - For single hidden layer: equivalent to taking the geometric average of all neural networks, with all possible binary masks
- Can be combined with unsupervised pre-training
- Beats regular backpropagation on many datasets
- **Ensemble:** Can be viewed as a geometric average of exponential number of networks.

Data Augmentation


- Increase the size of the training data
- It can be considered as a mandatory trick to improve predictions



[Slide from <https://www.analyticsvidhya.com/blog/2018/04/fundamentals-deep-learning-regularization-techniques/>]

Early Stop

- To select the number of epochs, stop training when validation set error increases (with some look ahead)



[Slide from <https://www.analyticsvidhya.com/blog/2018/04/fundamentals-deep-learning-regularization-techniques/>]

Batch Normalization

- Normalizing the inputs will speed up training (Lecun et al. 1998)
 - could normalization be useful at the level of the hidden layers?
- Batch normalization is an attempt to do that (Ioffe and Szegedy, 2015)
 - each unit's pre-activation is normalized (mean subtraction, stddev division)
 - during training, mean and stddev is computed for each minibatch
 - backpropagation takes into account the normalization
 - at test time, the global mean / stddev is used

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$



Learned linear transformation to adapt to non-linear activation function (γ and β are trained)

Batch Normalization

- Why normalize the pre-activation?
 - can help keep the pre-activation in a non-saturating regime (though the linear transform $y_i \leftarrow \gamma \hat{x}_i + \beta$ could cancel this effect)
- Use the **global mean and stddev** at test time.
 - removes the stochasticity of the mean and stddev
 - requires a final phase where, from the first to the last hidden layer
 - propagate all training data to that layer
 - compute and store the global mean and stddev of each unit
 - for early stopping, could use a running average

Computer Vision: Image Classification

- ImageNet LSVRC-2011 contest:
 - Dataset: 1.2 million labeled images, 1000 classes
 - Task: Given a new image, label it with the correct class

The screenshot shows the ImageNet website interface. At the top, there is a search bar with the placeholder "14,197,122 images, 21841 synsets indexed". Below the search bar are buttons for "SEARCH", "Home", "About", and "Explore Download". A message "Not logged in. Login | Signup" is displayed. On the left, a sidebar titled "Bird" provides a hierarchical list of synsets:

- marine animal, marine creature, sea animal, sea creature (1)
- scavenger (1)
- biped (0)
- predator, predatory animal (1)
- larva (49)
- acrodont (0)
- feeder (0)
- stunt (0)
- chordate (3087)
 - tunicate, urochordate, urochord (6)
 - cephalochordate (1)
 - vertebrate, craniate (3077)
 - mammal, mammalian (1169)
 - bird (871)
 - dickeybird, dickey-bird, dickybird, dicky-bird (0)
 - cock (1)
 - hen (0)
 - nester (0)
 - night bird (1)
 - bird of passage (0)
 - protoavis (0)
 - archaeopteryx, archeopteryx, Archaeopteryx lithographica (0)
 - Sinornis (0)
 - Ibero-mesornis (0)
 - archaeornis (0)
 - ratite, ratite bird, flightless bird (10)
 - carinate, carinate bird, flying bird (0)
 - passerine, passeriform bird (279)
 - nonpasserine bird (0)
 - bird of prey, raptor, raptorial bird (80)
 - gallinaceous bird, gallinacean (114)

On the right, there is a summary of statistics: "2126 pictures" and "92.85% Popularity Percentile". A "Wordnet IDs" icon is also present. Below these, there are three tabs: "Treemap Visualization", "Images of the Synset" (which is selected), and "Downloads". A grid of 4x6 images shows various bird species.


[Slide from Matt Gormley et al.]

Computer Vision: Image Classification

CNN for Image Classification

(Krizhevsky, Sutskever & Hinton, 2012)


15.3% error on ImageNet LSVRC-2012 contest



[Slide from Matt Gormley et al.]

CNNs for Image Recognition


○ Convolutional Neural Networks (CNNs)



[Slide from Matt Gormley et al.]

Convolutional Neural Network (CNN)

- Typical layers include:
 - Convolutional layer
 - Max-pooling layer
 - Fully-connected (Linear) layer
 - ReLU layer (or some other nonlinear activation function)
 - Softmax
- These can be arranged into arbitrarily deep topologies
- Architecture #1: LeNet-5



[Slide from Matt Gormley et al.]

What is a Convolution

- Basic idea:


- Pick a 3×3 matrix F of weights
- Slide this over an image and compute the “inner product” (similarity) of F and the corresponding field of the image, and replace the pixel in the center of the field with the output of the inner product operation

- Key point:

- Different convolutions extract different low-level “features” of an image
- All we need to vary to generate these different features is the weights of F

- A **convolution matrix** is used in image processing for tasks such as edge detection, blurring, sharpening, etc.

Ex: 1 input channel, 1 output channel



$$\begin{aligned}y_{11} &= \alpha_{11}x_{11} + \alpha_{12}x_{12} + \alpha_{21}x_{21} + \alpha_{22}x_{22} + \alpha_0 \\y_{12} &= \alpha_{11}x_{12} + \alpha_{12}x_{13} + \alpha_{21}x_{22} + \alpha_{22}x_{23} + \alpha_0 \\y_{21} &= \alpha_{11}x_{21} + \alpha_{12}x_{22} + \alpha_{21}x_{31} + \alpha_{22}x_{32} + \alpha_0 \\y_{22} &= \alpha_{11}x_{22} + \alpha_{12}x_{23} + \alpha_{21}x_{32} + \alpha_{22}x_{33} + \alpha_0\end{aligned}$$

[Slide from Matt Gormley et al.]

What is a Convolution

Input Image

0	0	0	0	0	0	0
0	1	1	1	1	1	0
0	1	0	0	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

Convolution

0	0	0
0	1	1
0	1	0

Convolved Image

3	2	2	3	1
2	0	2	1	0
2	2	1	0	0
3	1	0	0	0
1	0	0	0	0

0	1	0	0	0	0	0
0	1	1	1	1	1	0
0	0	1	1	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

0	1	0	0	0	0	0
0	1	1	1	1	1	0
0	0	1	1	1	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0

3	2			


[Slide from Matt Gormley et al.]

What is a Convolution



[Slide from Matt Gormley et al.]


What is a Convolution



[Slide from Matt Gormley et al.]

Downsampling by Averaging


- Suppose we use a convolution with stride 2
- Only 9 patches visited in input, so only 9 pixels in output



[Slide from Matt Gormley et al.]

Downsampling by Max-Pooling


- Max-pooling is another (common) form of downsampling
- Instead of averaging, we take the max value within the same range as the equivalently-sized convolution
- The example below uses a stride of 2



[Slide from Matt Gormley et al.]

CNN in protein-DNA binding

- Feature extractor for motifs




[Slide from Babak Alipanahi et al. 2015]

Recurrent Neural Networks

- Dataset for Supervised Part-of-Speech (POS) Tagging

Data: $\mathcal{D} = \{\mathbf{x}^{(n)}, \mathbf{y}^{(n)}\}_{n=1}^N$



[Slide from Matt Gormley et al.]

Recurrent Neural Networks

- Dataset for Supervised Handwriting Recognition


Data: $\mathcal{D} = \{\mathbf{x}^{(n)}, \mathbf{y}^{(n)}\}_{n=1}^N$



[Slide from Matt Gormley et al.]

Time Series Data

- **Question 1:** How could we apply the neural networks we've seen so far (which expect fixed size input/output) to a prediction task with variable length input/output?
- **Question 2:** How could we incorporate context (e.g. words to the left/right, or tags to the left/right) into our solution?



[Slide from Matt Gormley et al.]

Recurrent Neural Networks (RNNs)

inputs: $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units: $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$


outputs: $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity: \mathcal{H}

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$



[Slide from Matt Gormley et al.]

Recurrent Neural Networks (RNNs)

inputs: $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units: $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$


outputs: $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity: \mathcal{H}

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$



- If $T=1$, then we have a standard feed-forward **neural net with one hidden layer**
- All of the deep nets from last lecture required **fixed size inputs/outputs**


Recurrent Neural Networks (RNNs)

inputs: $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$
hidden units: $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$
outputs: $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$
nonlinearity: \mathcal{H}

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$
$$y_t = W_{hy}h_t + b_y$$

- By unrolling the RNN through time, we can **share parameters** and accommodate **arbitrary length** input/output pairs
- Applications: **time-series data** such as sentences, speech, stock-market, signal data, etc.



[Slide from Matt Gormley et al.]

Bidirectional RNN

inputs: $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units: $\overrightarrow{\mathbf{h}}$ and $\overleftarrow{\mathbf{h}}$

outputs: $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$


nonlinearity: \mathcal{H}

Recursive Definition:

$$\overrightarrow{\mathbf{h}}_t = \mathcal{H} \left(W_{x \overrightarrow{\mathbf{h}}} x_t + W_{\overrightarrow{\mathbf{h}} \overrightarrow{\mathbf{h}}} \overrightarrow{\mathbf{h}}_{t-1} + b_{\overrightarrow{\mathbf{h}}} \right)$$

$$\overleftarrow{\mathbf{h}}_t = \mathcal{H} \left(W_{x \overleftarrow{\mathbf{h}}} x_t + W_{\overleftarrow{\mathbf{h}} \overleftarrow{\mathbf{h}}} \overleftarrow{\mathbf{h}}_{t+1} + b_{\overleftarrow{\mathbf{h}}} \right)$$


$$y_t = W_{\overrightarrow{\mathbf{h}} y} \overrightarrow{\mathbf{h}}_t + W_{\overleftarrow{\mathbf{h}} y} \overleftarrow{\mathbf{h}}_t + b_y$$



[Slide from Matt Gormley et al.]

Deep Bidirectional RNNs

- Notice that the upper level hidden units have input from two previous layers (i.e. wider input)
- Likewise for the output layer




[Slide from Matt Gormley et al.]

Long Short-Term Memory (LSTM)

- Motivation:

- Vanishing gradient problem for Standard RNNs
- Figure shows sensitivity (darker = more sensitive) to the input at time t=1




[Slide from Matt Gormley et al.]

Long Short-Term Memory (LSTM)


- Motivation:

- LSTM units have a rich internal structure
- The various “gates” determine the propagation of information and can choose to “remember” or “forget” information



[Slide from Matt Gormley et al.]

Long Short-Term Memory (LSTM)



[Slide from Matt Gormley et al.]

Long Short-Term Memory (LSTM)

- **Input gate:** masks out the standard RNN inputs
- **Forget gate:** masks out the previous cell
- **Cell:** stores the input/forget mixture
- **Output gate:** masks out the values of the next hidden


$$i_t = \sigma (W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i)$$

$$f_t = \sigma (W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f)$$

$$c_t = f_t c_{t-1} + i_t \tanh (W_{xc}x_t + W_{hc}h_{t-1} + b_c)$$

$$o_t = \sigma (W_{xo}x_t + W_{ho}h_{t-1} + W_{co}c_t + b_o)$$


$$h_t = o_t \tanh(c_t)$$



[Slide from Matt Gormley et al.]

Deep Bidirectional LSTM (DBLSTM)

- How important is this particular architecture?
- Jozefowicz et al. (2015) evaluated 10,000 different LSTM-like architectures and found several variants that worked just as well on several tasks.



[Slide from Matt Gormley et al.]

Take home message

- Methods to prevent overfitting in deep learning
 - L2 & L1 regularization
 - Dropout
 - Data augmentation
 - Early stopping
 - Batch normalization
- CNN
 - Are used for all aspects of computer vision
 - Learn interpretable features at different levels of abstraction
 - Typically, consist of **convolution layers**, **pooling layers**, **nonlinearities**, and **fully connected** layers
- RNN
 - Applicable to sequential tasks
 - Learn context features for time series data
 - Vanishing gradients are still a problem – but LSTM units can help

References

- Matt Gormley. 10601 Introduction to Machine Learning:
<http://www.cs.cmu.edu/~mgormley/courses/10601/index.html>
- Barnabás Póczos, Maria-Florina Balcan, Russ Salakhutdinov. 10715 Advanced Introduction to Machine Learning:
<https://sites.google.com/site/10715advancedmlintro2017f/lectures>